

## Data Mining V: Preparing the Data

Computer Science 105  
Boston University

David G. Sullivan, Ph.D.

### The Data Mining Process

- Key steps:
  - assemble the data in the format needed for data mining
    - typically a text file
  - perform the data mining
  - interpret/evaluate the results
  - apply the results

## Denormalization

- Recall: in designing a database, we try to avoid redundancies by *normalizing* the data.
- As a result, the data for a given entity (e.g., a customer) may be:
  - spread over multiple tables
  - spread over multiple records within a given table
- To prepare for data warehousing and/or data mining, we often need to *denormalize* the data.
  - multiple records for a given entity → a single record
- Example: finding associations between courses students take.
  - our university database has three relevant relations: Student, Course, and Enrolled
  - we might need to combine data from all three to create the necessary training examples

## Transforming the Data

- We may also need to reformat or transform the data.
  - we can use a Python program to do the reformatting!
- One reason for transforming the data: many machine-learning algorithms can only handle certain types of data.
  - some algorithms only work with *nominal* attributes – attributes with a specified set of possible values
    - examples: {yes, no}
    - {strep throat, cold, allergy}
  - other algorithms only work with *numeric* attributes

## Discretizing Numeric Attributes

- We can turn a numeric attribute into a nominal/categorical one by using some sort of *discretization*.
- This involves dividing the range of possible values into subranges called *buckets* or *bins*.
  - example: an *age* attribute could be divided into these bins:  
child: 0-12  
teen: 12-17  
young: 18-35  
middle: 36-59  
senior: 60-

## Simple Discretization Methods

- What if we don't know which subranges make sense?
- *Equal-width binning* divides the range of possible values into N subranges of the same size.
  - $\text{bin width} = (\text{max value} - \text{min value}) / N$
  - example: if the observed values are all between 0-100, we could create 5 bins as follows:  
 $\text{width} = (100 - 0) / 5 = 20$   
bins: [0-20], (20-40], (40-60], (60-80], (80-100]  
[ or ] means the endpoint is included  
( or ) means the endpoint is not included
  - typically, the first and last bins are extended to allow for values outside the range of observed values  
(-infinity-20], (20-40], (40-60], (60-80], (80-infinity)
  - problems with this equal-width approach?

### Simple Discretization Methods (cont.)

- *Equal-frequency* or *equal-height binning* divides the range of possible values into N bins, each of which holds the same number of training instances.
  - example: let's say we have 10 training examples with the following values for the attribute that we're discretizing:  
5, 7, 12, 35, 65, 82, 84, 88, 90, 95  
to create 5 bins, we would divide up the range of values so that each bin holds 2 of the training examples:  
5, 7, 12, 35, 65, 82, 84, 88, 90, 95
- To select the boundary values for the bins, this method typically chooses a value halfway between the training examples on either side of the boundary.
  - examples:  $(7 + 12)/2 = 9.5$        $(35 + 65)/2 = 50$
- Problems with this approach?

### Other Discretization Methods

- Ideally, we'd like to come up with bins that capture distinctions that will be useful in data mining.
  - example: if we're discretizing *body temperature*, we'd like the discretization method to learn that 98.6 F is an important boundary value
  - more generally, we want to capture distinctions that will help us to learn to predict/estimate the class of an example
- Both equal-width and equal-frequency binning are considered *unsupervised* methods, because they don't take into account the class values of the training examples.
- There are *supervised* methods for discretization that attempt to take the class values into account.

## Discretization in Weka

- In Weka, you can discretize an attribute by applying the appropriate filter to it.
- After loading in the dataset in the *Preprocess* tab, click the *Choose* button in the *Filter* portion of the tab.
- For equal-width or equal-height, you choose the *Discretize* option in the *filters/unsupervised/attribute* folder.
  - by default, it uses equal-width binning
  - to use equal-frequency binning instead, click on the name of the filter and set the *useEqualFrequency* parameter to *True*
- For supervised discretization, choose the *Discretize* option in the *filters/supervised/attribute* folder.

## Nominal Attributes with Numeric Values

- Some attributes that use numeric values may actually be nominal attributes.
  - the attribute has a small number of possible values
  - there is no ordering to the values, and you would never perform mathematical operations on them
  - example: an attribute that uses numeric codes for medical diagnoses
    - 1 = Strep Throat, 2 = Cold, 3 = Allergy
- If you load into Weka a comma-separated-value file containing such an attribute, Weka will assume that it is numeric.
- To force Weka to treat an attribute with numeric values as nominal, use the *NumericToNominal* option in the *filters/unsupervised/attribute* folder.
  - click on the name of the filter, and enter the number(s) of the attributes you want to convert

## Removing Problematic Attributes

- Problematic attributes include:
  - irrelevant attributes: ones that don't help to predict the class
    - despite their irrelevance, the algorithm may erroneously include them in the model
  - attributes that cause overfitting
    - example: a unique identifier like *Patient ID*
  - redundant attributes: ones that offer basically the same information as another attribute
    - example: in many problems, date-of-birth and age provide the same information
    - some algorithms may end up giving the information from these attributes too much weight
- We can remove an attribute manually in Weka by clicking the checkbox next to the attribute in the *Preprocess* tab and then clicking the *Remove* button.

## Undoing Preprocess Actions

- In the Preprocess tab, the *Undo* button allows you to undo actions that you perform, including:
  - applying a filter to a dataset
  - manually removing one or more attributes
- If you apply two filters without using *Undo* in between the two, the second filter will be applied to the results of the first filter.
- *Undo* can be pressed multiple times to undo a sequence of actions.

## Dividing Up the Data File

- To allow us to validate the model(s) learned in data mining, we'll divide the examples into two files:
  - $n\%$  for training
  - $100 - n\%$  for testing: these should not be touched until you have finalized your model or models
  - possible splits:
    - 67/33
    - 80/20
    - 90/10
- You can use Weka to split the dataset for you after you perform whatever reformatting/editing is needed.
- If you discretize one or more attributes, you need to do so *before* you divide up the data file.
  - otherwise, the training and test sets will be incompatible

## Dividing Up the Data File (cont.)

- Here's one way to do it in Weka:
  - 1) shuffle the examples by choosing the *Randomize* filter from the *filters/unsupervised/instance* folder
  - 2) save the entire file of shuffled examples in Arff format.
  - 3) use the *RemovePercentage* filter from the same folder to remove some percentage of the examples
    - whatever percentage you're using for the training set
    - click on the name of the filter to set the percentage
  - 4) save the remaining examples in a new file
    - this will be our test data
  - 5) load the full file of shuffled examples back into Weka
  - 6) use *RemovePercentage* again with the same percentage as before, but set *invertSelection* to *True*
  - 7) save the remaining examples in a new file
    - this will be our training data